

The AI Hub and the Collection

A brief for leadership: what we built, what it can and cannot do, what it costs, and what we need from you.

For	Leadership; the platform, security and model-risk leads; the sponsors of the AI champions programme
Date	20 September 2026
Status	Reflects the repository at the date above; the readiness page in the repository is the live record

The short version

We took the working parts out of two internal AI projects, the first responder and the knowledge base, and turned them into a shelf of building blocks that any team can reuse: 29 of them, each tested, versioned and signed off by two people before it counts as ready. We put an employee AI hub in front of the shelf, implemented the hub's back end and an agent runtime so the whole thing runs inside the bank against our own identity provider, model service, secrets, Jira and Azure DevOps, and packaged it so it can be delivered without touching the public internet. Around it there's a programme, one engineer per team every two weeks, to turn the blocks into each team's own initiatives.


If you take one thing from this brief, take this: the agents in here can read, and they can draft, but they cannot change anything until a named person says yes to that exact change. That is not a policy we ask people to follow. It is how the code is built, and we can show you the tests.

The question you're probably asking

Most people's worry about AI agents is the right one: what stops it doing something we didn't ask for? Here is our answer, in the order the code enforces it.

An agent can only call the tools written in its template, and that list is signed. Every call goes through the same three checks, and the model cannot skip them. Reading is allowed. Writing is parked until the person who is using the agent confirms that exact action, once; a second confirmation of the same thing is refused. Anything bigger needs a second person. Anything involving money is refused outright. If the text the agent reads contains instructions, say a ticket someone has doctored, the session is marked and from then on it can only read. There is a kill switch. And every step is written to a record that cannot be edited without the edit showing.

One call through the loop. Reads run; a write waits for the acting person; a bigger change needs a second person; money is forbidden.




One call through the loop. The three hooks are ordinary code that runs the same way every time.

The employee assistant is simpler still. It reads pages from the knowledge base and tells you what they say, with the page cited. It has no tools. A claim it cannot back with a page is dropped or shown as unsupported.

What we've actually delivered

Twenty-nine components in six categories, every test and live example passing, each traced to the platform design specification with a test that says when the platform can replace it. The hub, pixel-identical to its design, with the sign-off queue and onboarding tracker added. The hub's API and an agent runtime, both on the Python standard library with no third-party dependency, both refusing to start in production with any fake left in. Containers, task definitions and least-privilege roles with placeholders where the bank's identifiers go. An offline release bundle that verifies itself. A user manual, a technical guide, a runbook, a security page and a readiness page.



How it sits inside the bank. Nothing on this diagram reaches the public internet; the hub's own fonts and scripts are served from inside.

What it costs

Two small containers and a storage volume each; that's the compute. Model usage is the variable cost, and it is bounded: every session has a budget of tokens, calls and time that the harness enforces, and the hub shows usage back per person and per cost centre. The bigger cost is people: an enablement lead, two to four hours a week per champion, and review time from AI security engineers, of whom we need at least two so nothing waits on one person. There's no licence and no external service.

What is not done, and who owns it

Everything that can be finished without the bank's accounts is finished. What remains is a list of things only the bank can supply, each with an owner in the readiness page: group ids and app registrations from identity; secrets, a signing key, an inference profile and a bucket from the platform and integrations teams; the model approved by model risk; the base image pinned to our approved one; our own listings loaded into the hub; and the sign-offs themselves, which we have deliberately not filled in, because a sign-off is a person's word and we won't fake one.

What we need from you

- 1 Name the enablement lead and two AI security engineers.
- 2 Approve a first, read-only production step: the hub in staging, and the agent reading real tickets and deploy histories but writing nothing. The first write is a separate decision, gated by a demonstration.
- 3 Ask model risk to review the inference profile and the first-read evaluation.
- 4 Confirm the champions programme's cadence and the first cohort.

The next ninety days

Weeks one to three: the open items closed and the staging deployment up; the first components used for real and signed. Weeks four to eight: the first cohort's briefs filed and two initiatives built on paved roads; the agent reading real incidents in staging. Weeks nine to thirteen: production for the hub and the read-only agent; the first write step proposed with its demonstration; the readiness page updated with measured numbers rather than estimates.

We'd rather you saw it than read about it. The walkthrough video is eight minutes, and a live demonstration with a poisoned ticket takes ten.