

The AI Hub and the Collection: leadership brief

What was built, why it matters, what it costs, what is governed, and what is needed to deploy it into the bank.

For	Leadership; the platform, security and model-risk leads; the sponsors of the AI champions programme
Date	20 September 2026
Status	Reflects the repository at the date above; the readiness page in the repository is the live record

1. In one paragraph

Two internal products, the first responder and the knowledge base, produced working AI controls, connectors and habits that every team building with AI needs. The Collection is how they spread: 29 self-contained building blocks in six categories, each versioned, tested, walked through and signed off by its owner and an AI security engineer, listed on the employee AI hub and published into the knowledge base. Around it, an AI champions programme (one engineer per team, every two weeks) turns those blocks into the team's own initiatives on approved paved roads. The hub, its API and an agent runtime are implemented and packaged to run inside the bank against its identity provider, model service, secrets, ticketing and deployment systems, with nothing fetched from the public internet.

2. What is delivered

Deliverable	State
29 components: 1 agent, 2 harnesses, 9 tools, 7 integrations, 3 patterns, 7 skills	Every test suite and live example green; each traced to the platform design specification with a replacement test
The employee AI hub	Pixel-identical to its design; configured at runtime; self-hosted assets; 43 tests and a pixel guard at zero
The hub's API behind the bank's identity provider	Implemented on the standard library; 17 tests including the OIDC path and the assistant backends; a browser flow proven end to end
An agent runtime	The incident first-read agent over MCP and a run API, wired to Jira, Azure DevOps, Bedrock, KMS and the identity provider by configuration
Governance surfaces	Sign-off queue and form, onboarding tracker, categories; the manifest is the record and the commit the signature
Packaging	Containers with fail-closed entrypoints, ECS task definitions, least-privilege roles; an offline release bundle; one verify script for GitHub and Azure Pipelines
Operating pages	Production readiness, runbook, security notes, handover, configuration guide, contributing contract; a user manual and a technical guide

3. Why it matters

- **Speed with control.** A team starts from a signed block with a five-minute README instead of from a blank page, and every block enforces the same rules: a model only cites, a write waits for a person, a tainted input caps the session at reads, nothing holds a credential.
- **One front door.** Employees see what they may use and ask for the rest; leaders see what is in flight and where each initiative stands.
- **A path to the platform.** Each block names what in the platform design replaces it and the test that proves the replacement; the work is not throwaway and not a second platform.
- **People, not just software.** The champions programme, the sign-off roles and the onboarding stages make adoption a practice with names on it.

4. What it costs to run

Item	Driver	Note
Compute	Two small containers (hub, agent) on Fargate, a record volume each	Standard library only; no licence; no third-party service
Model usage	Assistant turns and agent first reads on Bedrock	Budgets per session are enforced; the hub shows usage back; charge-back per cost centre is in the design
People	The enablement lead; two to four hours a week per champion; AI security engineers' review time	Two AI security engineers at minimum so no component waits on one person
Platform	Identity registrations, secrets, a KMS key, an inference profile, a bucket	One-time setup, named in the readiness page

5. How it is governed

Two sign-offs at every version, attested on a form the server enforces; fakes that cannot reach production; credentials never held; a hash-chained record verified at every start and exported nightly; logs that carry identifiers only; a threat model mapped to tests. Model risk gets the engine, stage and token counts on the chain and a feedback question on every answer. The categories say what a thing is and who may use it: a skill anyone, an agent only through its harness.

6. What is open, and who owns it

Open item	Owner
Group ids, app registrations, the groups claim on tokens	Identity engineer
Secrets, the KMS key, the inference profile, the bucket, the task roles applied	Integrations engineer, platform team
The model approved and the first-read judge validated	Model risk
The base image pinned and images built in the bank's registry	Integrations engineer
The bank's own listings and the identity map filled in	Platform team
The first real use of each component and the sign-offs themselves	Component owners, AI security engineers
The pages published into the bank's knowledge base	Enablement lead
The MCP conformance suite run with a reviewed expected-failures file	Platform team

None of these is a design gap: each is an account, an id, a key or a person. When they are closed, deployment is a configuration change on the build that already passed every gate.

7. Decisions needed

- 1 Name the two AI security engineers and the enablement lead.
- 2 Approve the first production increment: the hub in staging with the platform runtime's assistant (or none), and the agent reading real systems and writing to none.
- 3 Approve the inference profile with model risk and the charge-back rule for model usage.
- 4 Confirm the champions programme's cadence and the first cohort, one engineer per team.

8. The next ninety days

- Weeks 1 to 3: the open items closed; staging deployment; the first components used once for real and signed.
- Weeks 4 to 8: the first cohort's briefs filed; two initiatives built on paved roads from the collection; the first agent reading real incidents in staging.
- Weeks 9 to 13: production for the hub and the read-only agent; the first write increment gated by its demonstration; the readiness page re-baselined with measured numbers.