

The AI Hub and the Collection

A user manual for employees and AI champions: what is there, how to use it, and who to ask.

For	Everyone who signs in to the hub; AI champions; component owners; AI security engineers
Date	20 September 2026
Status	Reflects the repository at the date above; the readiness page in the repository is the live record

1. What this is, in plain words

The **AI Hub** is the front door to the company's AI services. You sign in with your normal work account, see what you may use, open an assistant, and ask for access to the rest. Behind it is the **Collection**: a shelf of ready-made AI building blocks that engineers reuse instead of rebuilding, each one checked, versioned and signed off by two people before it is called ready.

You do not need to know how any of it works to use it. This manual walks through the screens in the order you will meet them, then explains the two things that involve people beyond just using it: proposing a new AI use for your team, and signing a building block off if you are its owner or an AI security engineer.

Three rules the hub keeps for you. An assistant only answers from sources it can cite, and shows them. Anything that would change a system waits for a person to confirm it. Nothing you type is used to change the assistant's instructions; text is treated as a question, never as a command.

2. Signing in

Open the hub's address in your browser and sign in with your company account; the identity provider does the work and the hub never sees your password. What you see afterwards depends on your role and team: the hub shows only what you may use, and lets you ask for the rest. If a screen you expected is missing, that is access, not a fault; the request button is on the page.

3. Discover: what you may use today

Discover is the catalog. The top row is what is already open to you. Below it, tabs group everything else: assistants, agents, knowledge, tools, and the paved roads (the approved ways to build something new). The search box at the top finds any of them by name or by what they do.

The screenshot shows the CrossRiver AI Hub interface. At the top, there's a navigation bar with 'Discover', 'My workspace', 'Build', and 'Learn'. A search bar is on the right. Below the navigation, a section titled 'What do you want to do with AI today?' features a search input and several suggested queries. The main area is divided into two sections: 'Available to you' and 'Browse the catalog'. The 'Available to you' section displays four cards for different AI agents: 'Employee assistant', 'Investigation triage', 'Policy and procedures', and 'Compliance narration'. Each card includes a description, access levels (e.g., R2, R4), and an 'Open' button. The 'Browse the catalog' section shows a list of all agents, with the 'Agents' tab selected, displaying four more cards: 'Investigation triage', 'First responder', 'Compliance narration', and 'Payments exception agent'.

Discover with the Agents tab open. Each card names the service, its road, whether it is generally available or in preview, and how you reach it.

A card says one of four things about access: **Open** (click to use it), **Request access** (one click sends the request to your lead), **View** (you can read about it, your role does not use it), or **Contract** (a service another team runs; the page says who to ask).

4. A listing page

Every service and every building block has a page. Read it top to bottom: what it does, what it can read and do (each operation with its tier, from read-only to changes that need confirmation), the evidence behind it (tests, a step-by-step walkthrough, a live example, the knowledge base page), the owner and how to get support, and its versions.

Incident first read agent
An agent, complete: a template (role, stages, tools by tier, never list), tools reached only through the harness, a cited first read, one W1 proposal to confirm

road R1 python agent - ready spec \$5.1, \$5.2, \$5.5, \$4.8, \$4.5 agents security governance

Open in the portal → Try in playground

Template
TEMPLAT 2
role: the first read of an incident: an assistant that reads the t

Tools
cited-llm-engine, audit-chain

Harness
governed-action-loop
ladder L2

Language
python
standard library only

Status
ready
snapshot 2026-09-19

Pairs with
4
governed-action-loop, cited-llm-engine, teams-graph-connector, comm-templates

What it does

- ✓ The first thing a responder wants when an alert fires is a first read: what changed, what is failing, the leading hypothesis, with every claim pointing at a source, and at most one proposed action they can confirm or ignore. This is that agent, small enough to read in ten minutes, and the reference for what an agent in the collection is: a **template** that says what it is and may call, **tools** it reaches only through the harness, and a **harness** that enforces tiers, taint, budgets and the record. Copy it to start your own agent; replace the template and the fakes, keep the loop.
- ✓ The catalog the harness loads is built and signed from the template's `tools`; a call to anything else fails at the catalog hook.
- ✓ A write (W1) parks with a hash of the exact tool and arguments; only the acting person confirms; the confirmation is consumed once.
- ✓ A proposal is refused before any model call when the context is tainted, and the harness caps a tainted session at reads.
- ✓ The model sees projected, masked, fenced sources only; claims without a citation are dropped; a malformed answer is refused.

Sign-off v1.0.0 **owner pending; AI security pending**

Owner	pending	pending
AI security	pending	pending
Stage	built - next: use it once in a real project and record where (used_in, or the owner's sign-off form)	

Owner sign-off on Incident first read agent 1.0.0

- ☐ I ran the tests and they are green python3 -m unittest discover -s tests -t . && python3 example.py
- ☐ I ran the live example python3 example.py
- ☐ I read the walkthrough end to end WALKTHROUGH.md

A listing page. The right column carries the sign-off card, the getting-started steps, the owner and the versions.

The chip row under the title tells you at a glance: the road, the language, whether it is ready, which parts of the platform design it implements, and its tags. A component that both its owner and an AI security engineer have signed is marked **signed**; otherwise the hub shows it as a preview.

5. The employee assistant

Open the assistant from Discover or from your workspace, type a question, and read the answer. Three things are worth knowing:

- **Citations.** Claims in an answer are marked, and the sources appear under it with their classification. A claim without a source is shown as unsupported; treat it as a hint, not a fact.
- **Tool calls.** When the assistant looks something up, the step appears in the conversation with what it searched for. You can see what it did, not only what it said.
- **Feedback and handoff.** Each answer ends with one question: did this answer yours? Answer it; it is how the assistant is measured. If you would rather talk to a person, use the handoff and the conversation is routed to one.

If the assistant answers that a source looked like an instruction rather than evidence and it did not use it, that is the protection working: something in a page or a ticket tried to steer it. Rephrase, or open the page directly.

6. Your workspace

My workspace gathers what is yours: the assistants you use, your team's services and where each is in its journey to production, your pending requests, your usage, and the sandbox playground for trying things. Settings (theme, density, language, accessibility, notifications) are under the same menu; they

apply at once and travel with your account.

7. Proposing an AI use for your team: the brief

Build opens the intake brief: one page in six short sections, saved as you go. Fill it in the order it asks:

- 1 **Use case.** A name, what happens today, the channel (your team, customers, partners, or a batch), and the team.
- 2 **People.** The business owner, the product owner, and the domain expert who will label examples, with the hours a week they can give.
- 3 **Data and tools.** The systems of record it reads, the tools it may call and their tiers, the data classes, and the ceiling on what it may do.
- 4 **Model.** What the model needs to be, the classification ceiling, and whether a substitute is acceptable.
- 5 **Outcome.** One metric, its unit, today's figure, the target, and when the baseline was measured.
- 6 **Review.** What happens next, acknowledged.

The brief estimates the monthly model spend and the review hours as you fill it, and recommends a paved road. Filing it puts it in front of the platform lead; a brief that would let an AI change systems is filed by your team lead, and the form tells you so.

8. For AI champions: the programme in one page

One engineer per team, every two weeks, forty-five minutes: ten minutes of what shipped, one deep dive, and fifteen minutes on the collection. Each champion carries one initiative at a time, written as a brief, built from the collection on a paved road, and brought back as something everyone can reuse. Your first week: sign in, run one live example from the Tools tab, read one practice page in the knowledge base, and draft your brief.

The knowledge base holds the practices behind the components, the paved roads, the skills anyone can follow, and the programme's own pages; the Learn tab links to it.

9. For owners and AI security engineers: signing a component off

A building block reaches the shelf when two people have signed it at its current version: its **owner** (named on the component, signs by name) and an **AI security engineer** (signs by role). Signing means attesting four things on the form: you ran the tests and they are green, you ran the live example, you read the walkthrough end to end, and you read the rules it enforces and its known limits. The owner also names the project it was first used in for real.

DiscoverMy workspaceBuildLearn

Build / Sign-offs

Component sign-offs

A component reaches the shelf when its owner and an AI security engineer have each signed it at its current version. Sign here; the manifest in the repository is the record and the commit is the signature.

0 signed

23 awaiting owner

23 awaiting AI security

The queue

All 23

Open

Yours to sign 23

Recorded here

COMPONENT	VERSION	STAGE	OWNER	AI SECURITY	
Audit chain tool · python · owner gil.klainert	1.0.0	built	pending	pending	Close

Owner sign-off on Audit chain 1.0.0

☐ I ran the tests and they are green python3 -m unittest discover -s tests -t . -v
 ☐ I ran the live example python3 example.py
 ☐ I read the walkthrough end to end WALKTHROUGH.md
 ☐ I read the rules it enforces and its known limits README.md

Used once for real: which project

The owner signs after one real use. Name the project; it is recorded in the manifest.

Note (optional)

Where it ran and what you changed to make it fit.

recorded as 1.0.0; the commit that follows is the signature

Comm templates
tool · python · owner gil.klainert

1.0.0

built

pending

pending

Sign

Governed action loop
tool · python · owner gil.klainert

1.0.0

built

pending

pending

Sign

Ids-only logging
tool · python · owner gil.klainert

1.0.0

built

pending

pending

Sign

You

You may sign as **owner** on the components that name you.

23 waiting on you. Before you sign: run the tests and the live example, read the walkthrough and the rules. The form asks you to attest to each.

Into the repository

0 recorded here

A sign-off recorded here is a request until it is in the manifest. Export the queue, apply it, and commit; the shelf tool re-runs the tests before it writes anything.


```
python3 tools/shelf.py --apply-signoffs shelf-signoffs.json
python3 tools/shelf.py --write
git commit -am "Sign off <component> <version>"
```

The onboarding process

The sign-off queue under Build. The form opens in the row; the right column says what you may sign and how a recorded sign-off becomes a record.

The queue lists every component with its version and both sign-offs; the filter **Yours to sign** shows what waits on you. A sign-off recorded here is a request until an engineer applies it to the component and commits; the queue's download is that step. A new version of a component asks both people to sign again.

Internal. Placeholders in capitals are not real identifiers. Nothing here is a credential.



The onboarding tracker: each component's stage on its way to the shelf, and the checklists for a new champion, an owner, and an AI security engineer.

10. Getting help

Question	Where
I cannot see a service I need	The request button on its card or page; your lead approves
The assistant's answer looks wrong	Answer the feedback question with no; open the cited page; use the handoff for a person
I want to propose something	Build: start a brief; the platform lead reviews it
I am a champion and stuck	The champions channel and the office hour named in the hub's footer
A component I own needs a change	Its owner page names the process; the technical guide has the details
Something looks like a security problem	The AI security engineers who sign components off; do not wait for the next meeting